

BRENDA ONG

THREAT INTELLIGENCE · INFLUENCE OPERATIONS & OSINT · COUNTER-DISINFORMATION LEADERSHIP

Coordinated-Inauthentic-Behavior Detection · Threat-Actor Analysis · Cross-Platform Investigation · LLM Adversarial Evaluation

Singapore · Open to relocation globally · +65 9737 9509 · bihuiong@outlook.com

brendaong.com · linkedin.com/in/brenda-ong-4655aa3a

SUMMARY

Threat-intelligence and counter-disinformation leader with 13+ years detecting, investigating and disrupting influence operations, coordinated inauthentic behavior and online abuse across law enforcement, INTERPOL and frontier technology. Built and led a national 24/7 OSINT and investigation-triage function from scratch, trained 28-person team running 100+ cases a day at a 99% SLA that disrupted 20+ coordinated influence and smear campaigns. Set the casework quality bar, cut detection time over 50%, and briefed government and executive stakeholders. Hands-on with SQL, Python and LLM adversarial evaluation; MIT Sloan MBA; native Mandarin Chinese with geopolitical fluency.

CORE COMPETENCIES

Influence Operations & Coordinated Inauthentic Behavior · OSINT Tradecraft · Counter-Disinformation · Threat-Actor Attribution · Cross-Platform Campaign Analysis · Coordinated-Network Detection · Abuse Signals & Behavioral Clustering · LLM Weaponization & Adversarial Evaluation · Investigative & Threat-Intel Team Leadership · Government & Executive Intelligence Sharing · SQL, Python & Tableau · Mandarin Chinese

PROFESSIONAL EXPERIENCE

Principal Scientist, Disruptive Technologies Office | Home Team Science & Technology Agency Singapore

Jul 2025–Present

Lead responsible-AI programs for frontline agencies; hands-on LLM evaluation, adversarial testing and narrative-analysis tooling.

- Prototyped public-sentiment forecasting tool for policy announcements, modelling how narratives spread and shift across audiences.
- Built an LLM behavior-drift dashboard under adversarial prompting to measure model steering, jailbreak susceptibility and mitigation effectiveness.
- Built an AI-assisted analyst workflow with Claude Code and Codex to triage high-priority satellite imagery for human review, cutting report-prep time 30%.
- Translated LLM failure modes into evaluation thresholds, human-review requirements and escalation criteria for deployment.

Co-Founder & Chief Operating Officer | LymeAlert Boston, MA

Jun 2023–Jul 2025

Part-time since Jul 2025. Built the operating and safety-escalation systems of an early-stage diagnostics startup.

- Cut manual reporting and coordination ~10 hours/week (66.7%) by automating reporting and issue-tracking across five workstreams.
- Established a product-complaint and safety-escalation workflow, improving milestone completion from ~65% to 90%.

Assistant Director, Public Communications (Threat Intel, OSINT & Counter-Disinformation) | Singapore Police Force

Singapore

Feb 2021–May 2023

Official police spokesperson. Owned a national 24/7 OSINT, influence-operations detection, enforcement-referral and crisis operation.

- Detected and disrupted 20+ coordinated inauthentic influence and smear campaigns against public institutions via OSINT verification and cross-government cells.
- Ran a national 24/7 intelligence and triage operation at 100+ incidents a day, sustaining 99% response SLA across a 28-person team.
- Cut misinformation mean-time-to-detection by 50%+ to 3hrs using OSINT, social-listening and rules-based alerts across major platforms.
- Traced 3 coordinated campaigns across Facebook and TikTok amplified via Telegram, linking on-platform activity to broader ops.
- Improved human/automated detection agreement for 20%+ more cases yearly by analyzing false positives and converting edge cases into evaluation tests.
- Set casework quality bar for a 28-person function through 500+ monthly decision reviews, calibration sessions and targeted retraining.
- Attributed 5 campaigns to coordinated groups/foreign-state actors and 6 scam operations to criminal groups using OSINT and social network analysis.

Head of Operations, Tanglin Division | Singapore Police Force

Aug 2019–Feb 2021

Led operational readiness and incident response for a command of ~1,500 officers.

- Scaled a counter-terrorism rapid-response capability to national adoption by building its operating model, activation rules, training and QA cadence.
- Commanded 30+ major security operations using Tableau hotspot/capacity dashboards, and ran adversarial red-teaming with foreign security agencies.

Deputy Team Leader, Crisis Negotiation Unit (concurrent) | Singapore Police Force Jul 2017–May 2023

- Stabilized 40+ threat-to-life incidents using behavioral risk assessment and structured de-escalation; authored escalation protocols still in use.

Commanding Officer, Bukit Panjang NPC | Singapore Police Force Jan 2018–Aug 2019

- Built a Python/Tableau demand model redesigning deployment, cutting five major offences 30% and earning a top-five national rating.
- Co-authored an AI-ethics policy paper feeding Singapore’s national AI Governance Framework and Advisory Council.
- Lifted team engagement from 80% to 93% (pulse survey) by redesigning team structure, coaching and development pathways.

Operations Officer, Future Operations Division | Singapore Police Force Jan 2015–Jan 2018

Founded a four-person futures and technology-strategy unit for national public safety.

- Cut duplicative systems 42.9% across 10 departments by leading a 10-year joint operations-investigations-intelligence technology masterplan.
- Raised pilot-to-adoption conversion from 10% to 50% by adding stage gates, success measures and post-pilot ownership.

Executive, Strategic Planning Directorate | INTERPOL General Secretariat Lyon, France Jun 2014–Jan 2015

- Produced INTERPOL’s Environmental Scan series, forecasting emerging security threats across a 195-country membership for senior-leader strategy.

Investigation Officer, Central Police Division | Singapore Police Force Singapore Sep 2012–Jun 2014

- Solved 20+ investigations a month at 85%+ solvability, reviewed cybercrime casework, and built the investigative rulebook and case-exhibit tracking system.

EDUCATION

MBA, Sloan Fellows Program - MIT Sloan School of Management	2024
<i>ASEAN Fellow · Specialisation in Business Analytics; coursework in AI & Machine Learning · GPA 4.9/5.0</i>	
MA, Security Studies - Georgetown University, School of Foreign Service	2012
<i>Outstanding Academic Achievement · Concentration: International Security · GPA 3.89/4.0</i>	
BA (Hons, First Class), Politics, Psychology & Sociology - University of Cambridge	2011
<i>Kerslake Scholar, Emmanuel College; Awarded Honorary MA in 2015</i>	
MTech in AI Systems (in progress) - National University of Singapore, ISS	2026 – present

TOOLS, CREDENTIALS & LANGUAGES

Technical & data: SQL · Python · R · Tableau · Power BI · Excel · Claude Code · Codex · LLM evaluation & adversarial prompting · OSINT & social-listening platforms · Social network analysis

Credentials: AI Security, Lakera (prompt-injection defence) · Data Science Professional, DataCamp · AI Nanodegree, Heicoders · Certified ScrumMaster · Trained Crisis Negotiator

Recognition: The Commendation Medal (COVID-19), National Award, Prime Minister’s Office, Singapore, 2023

Languages: English (native) · Mandarin Chinese (native / fluent) · French (basic)

SELECTED PROJECTS

Four open research projects on AI misuse, safety incidents and enforcement operations. Code, data and reports at brendaong.com and github.com/bibiong.

AI Misuse Threat Intelligence Atlas — Threat-intelligence corpus. Eight frontier-lab and platform threat reports coded into 55 individually cited operation records against a purpose-built kill chain and a verified MITRE ATLAS mapping.

- Found cross-language lure localisation to be the most converged technique in the public record (20 of 55 operations, 5 of 7 actor types, all three contributing publishers), with no corresponding MITRE ATLAS or ATT&CK technique; measured ATLAS coverage of documented misuse at 25% of 252 coded technique instances.

AI Incident After-Action Library — Twelve structured after-action reviews of public AI safety incidents — timeline, contributing factors, detection gap, corrective actions with owners — with schema validation and generated cross-library analysis.

- Found 11 of 12 failures were first surfaced by someone outside the operating organisation; tracked 66 corrective actions across the library, of which 50% are externally verifiable.

Policy-to-Eval Harness — Builds the missing step between a written usage policy and a decision an evaluator can score. 400 borderline prompts across 5 open-weight models, 2,000 scored responses.

- Measured leakage at 1.8x over-refusal (14.2% of prohibited requests answered against 7.7% of permitted requests refused), inverting the project's framing; reported judge calibration (Cohen's kappa 0.68 weighted) and the rubric defect producing systematic one-directional disagreement.

Enforcement Operations Simulator — Discrete-event simulation of a 1,700-alert/day Trust & Safety enforcement queue (~4,100 lines of Python, 82 tests), plus the written operating model it exists to test.

- Showed the aggregate SLA hides band-level failure (98.6% aggregate while missing P1 at 93.1% against a 95% target), and that stratified QA sampling detects single-reviewer policy drift at 5% sampling where uniform sampling fails at 50%.